

MIT Doctoral Dissertation  
PUBLISH ABSTRACT ONLY  
INFORMATION\_\_\_\_\_  
Abstract No.***DO NOT WRITE IN THIS SPACE***

Vol/Issue \_\_\_\_\_

School Code \_\_\_\_\_

Advisor \_\_\_\_\_

***PLEASE TYPE OR PRINT***

## PERSONAL DATA

1. Full name (as it appears on dissertation title page)

\_\_\_\_\_  
(last) (first) (middle)

2. Year of birth (optional) \_\_\_\_\_

3. Present mailing address \_\_\_\_\_

\_\_\_\_\_

Future mailing address \_\_\_\_\_

\_\_\_\_\_

Effective date of future mailing address \_\_\_\_\_

Home telephone \_\_\_\_\_ Business telephone \_\_\_\_\_

## DOCTORAL DEGREE DATA

4. Full name of university conferring degree
- Massachusetts Institute of Technology

5. Degree awarded (check one)      Ph.D.      Sc.D.

6. Year degree awarded \_\_\_\_\_

7. IMPORTANT: Attach a copy of your dissertation title page and abstract to this form. Please be certain that the name of your dissertation supervisor is included on both.

8. Subject categories for your dissertation. Enter 4-digit code from 'Subject Categories' list found on the opposite side of this form, and write in the category selected. You may enter two additional codes and categories on the lines provided.

Code \_\_\_\_\_ Category \_\_\_\_\_

Code \_\_\_\_\_ Category \_\_\_\_\_

Code \_\_\_\_\_ Category \_\_\_\_\_

(Optional) List up to five additional words from your dissertation not already found in *either* your title *or* abstract which would be useful for database access.

a. \_\_\_\_\_ b. \_\_\_\_\_ c. \_\_\_\_\_

d. \_\_\_\_\_ e. \_\_\_\_\_



[**TODO:** Make sure to exclude no files (not rewriting)]



# Performance Engineering of Proof-Based Software Systems at Scale

by

Jason S. Gross

Submitted to the Department of Electrical Engineering and Computer  
Science

in partial fulfillment of the requirements for the degree of  
Doctor of Philosophy in Computer Science and Engineering  
at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

June 2021

© Jason S. Gross, MMXXI. All rights reserved.

The author hereby grants to MIT permission to reproduce and to  
distribute publicly paper and electronic copies of this thesis document  
in whole or in part in any medium now known or hereafter created.

Author .....  
Department of Electrical Engineering and Computer Science  
August 19, 2015

Certified by .....  
Adam Chlipala  
Associate Professor of Computer Science  
Thesis Supervisor

Accepted by .....  
Leslie A. Kolodziejski  
Chair, Department Committee on Graduate Students



# Performance Engineering of Proof-Based Software Systems at Scale

by  
Jason S. Gross

Submitted to the Department of Electrical Engineering and Computer Science  
on August 19, 2015, in partial fulfillment of the  
requirements for the degree of  
Doctor of Philosophy in Computer Science and Engineering

## Abstract

Formal verification is increasingly valuable as our world comes to rely more on software for critical infrastructure. A significant and under-studied cost of developing verification, especially at scale, is the computer performance of proof generation.

This thesis aims to be a partial guide to identifying and resolving performance bottlenecks in dependently-typed tactic-driven proof assistants such as Coq.

We present a survey of the landscape of performance issues in Coq, with a number of micro- and macro-benchmarks. We describe various metrics that allow prediction of slowness, such as term size, goal size, and number of binders, and note the occasional surprise-lack-of-bottleneck for some factors, such as total proof term size. To our knowledge such a roadmap to performance bottlenecks is a new contribution of this thesis.

We identify three main categories of workarounds and partial solutions to slowness: design of APIs of Gallina libraries; changes to Coq’s type theory, implementation, or tooling; and automation design patterns, including proof by reflection. We present lessons drawn from the case-studies of a category-theory library, a proof-producing parser-generator, and a verified compiler and code generator for low-level cryptographic primitives.

The central new technical contribution presented by this thesis is a reflective framework for partial evaluation and rewriting, already used to compile a code-generator for field arithmetic cryptographic primitives which generates code currently used in Google Chrome. We believe this prototype is the first scalably performant realization of an approach for code specialization which does not require adding to the trusted code base. Our extensible engine, which combines the traditional concepts of tailored term reduction and automatic rewriting from hint databases with on-the-fly generation of inductive codes for constants, is also of interest to replace these ingredients in proof assistants’ proof checkers and tactic engines. Additionally, we use the development of this framework itself as a case-study for the various performance issues that can arise when designing large proof libraries.

Finally, we present a novel method of simple and fast reification, developed and published during the course of doctoral study.

Thesis Supervisor: Adam Chlipala

Title: Associate Professor of Computer Science